

Building an Agentic RAG Model for Crypto Trading with Haystack, Weaviate, and OpenRouter.ai

Grok 3, xAI

July 7, 2025

Contents

1	Introduction	2
2	System Architecture	2
2.1	Architecture Components	2
2.2	Architecture Diagram (Text Description)	3
2.3	Robustness Recommendations	3
3	Trading Strategy	3
3.1	Indicator Categories	3
3.2	Confluence Strategy	4
4	Setting Up Weaviate Document Stores	4
4.1	Weaviate Installation and Configuration	5
4.2	Vector Database 1: CoinMarketCap API Data	5
4.3	Vector Database 2: Binance API Data with Technical Analysis	6
4.4	Vector Database 3: News and Social Media Sentiment	8
4.5	Robustness Recommendations	9
5	Building the Haystack RAG Pipeline	9
5.1	Haystack Components	9
5.2	Pipeline Setup	10
6	Developing the Haystack Agent	10
6.1	Agent Setup	11
7	Building the Frontend Application	12
8	Integration and Testing	13
9	Deployment	13
9.1	Deployment Setup	13
9.2	Deployment Steps	13
9.3	Maintenance	14
10	Advanced Features	14
11	Regulatory Compliance	14
12	Conclusion	15

1 Introduction

This guide provides a detailed roadmap for building an Agentic Retrieval-Augmented Generation (RAG) model for cryptocurrency trading using the **Haystack framework** (v2.0+), **Weaviate** as the document store, and **OpenRouter.ai** for multi-model LLM integration. The system leverages three vector databases to store data from CoinMarketCap APIs, Binance APIs with technical analysis, and news/social media platforms (e.g., X) with sentiment analysis. A robust trading strategy combining multiple technical indicators and risk management principles is integrated to maximize profit and minimize risk. The system empowers a frontend application for users to create intelligent AI trading agents that make data-driven decisions and execute trades on selected cryptocurrency pairs.

The system supports:

- Real-time market data analysis with technical indicators.
- Sentiment-driven trading strategies.
- Autonomous trade execution with user-defined parameters.
- A user-friendly React-based frontend for agent configuration and monitoring.

2 System Architecture

The architecture integrates Haystack's RAG pipeline, Weaviate for vector storage, OpenRouter.ai for generative AI, and a robust trading strategy.

2.1 Architecture Components

- **Data Sources:**
 - **Vector Database 1 (CoinMarketCap):** Stores market data (price, market cap, volume).
 - **Vector Database 2 (Binance):** Stores trading data with technical indicators (e.g., EMA, RSI, MACD, Bollinger Bands).
 - **Vector Database 3 (Sentiment):** Stores news and X posts with sentiment scores.
- **Haystack RAG Pipeline:**
 - **Document Store:** Weaviate with three classes for data storage.
 - **Embedder:** sentence-transformers/all-MiniLM-L6-v2.
 - **Retriever:** WeaviateHybridRetriever for BM25 and dense search.
 - **Generator:** OpenRouter.ai's LLM (e.g., LLaMA 3.1 70B).
 - **Pipeline:** Combines components for query processing.
- **Haystack Agent:** Autonomous agents using tools for data retrieval, indicator computation, and trade execution.
- **Frontend Application:** React with Tailwind CSS and Chart.js for visualization.
- **Execution Layer:** Binance API for trade execution with safeguards.

2.2 Architecture Diagram (Text Description)

- **User Input:** Users configure agents via the frontend (e.g., BTC/USDT, 2% risk).
- **Haystack Agent:** Queries the RAG pipeline and executes trades based on a confluence strategy.
- **RAG Pipeline:**
 - **Query Processing:** Parses user goals (e.g., “Find oversold BTC/USDT with positive sentiment”).
 - **Retriever:** Fetches data from Weaviate.
 - **Generator:** Synthesizes trading signals using OpenRouter.ai.
- **Weaviate Document Store:** Stores embeddings for all data sources.
- **Output:** Trading signals displayed on the frontend or executed via Binance API.

2.3 Robustness Recommendations

- Use a **microservices architecture** for scalability.
- Integrate **CoinGecko** as a backup for CoinMarketCap.
- Use **Kafka** for high-frequency data ingestion.
- Deploy on **AWS** with auto-scaling.

3 Trading Strategy

The trading strategy combines multiple technical indicators and risk management to maximize profit and reduce risk, emphasizing confluence for high-confidence decisions.

3.1 Indicator Categories

- **Trend Indicators:**
 - **EMA/SMA:** 50 EMA > 200 EMA for bullish trend.
 - **MACD:** Bullish/bearish crossovers for momentum shifts.
 - **Parabolic SAR:** Trailing stops and reversal points.
- **Momentum Indicators:**
 - **RSI:** Overbought (>70), oversold (<30).
 - **Stochastic Oscillator:** Short-term reversal signals.
 - **CCI:** Price deviations from historical norms.
- **Volatility Indicators:**
 - **Bollinger Bands:** Detect squeezes and breakouts.
 - **ATR:** Guide stop-loss levels based on volatility.
- **Volume Indicators:**
 - **OBV:** Confirm trend strength.

- **VWAP:** Identify institutional interest.
- **Support & Resistance:**
 - **Horizontal S/R:** Historical reaction levels.
 - **Fibonacci Retracement:** 0.618 level for pullbacks.
 - **Pivot Points:** Intraday decision-making.
- **Price Action & Patterns:**
 - **Candlestick Patterns:** Engulfing, pin bars for reversals.
 - **Chart Patterns:** Triangles, head and shoulders for breakouts.
- **Risk Management:**
 - **Position Sizing:** Limit to 1–2% of account per trade.
 - **Trailing Stop-Loss:** Lock in gains.
 - **Risk-to-Reward Ratio:** Minimum 1:2.

3.2 Confluence Strategy

Buy Setup:

- 50 EMA > 200 EMA (bullish trend).
- RSI crosses above 30 (momentum shift).
- Price bounces from 0.618 Fibonacci support.
- MACD bullish crossover.
- Volume spike (OBV increasing).

Sell/Exit Setup:

- RSI > 70 with bearish divergence.
- MACD bearish crossover.
- Price hits resistance with extended ATR.
- Parabolic SAR flips above price.

Risk Management:

- Stop-loss: Set at 2x ATR below entry.
- Take-profit: Target 2x risk or next Fibonacci level.
- Position size: 1% of account to limit exposure.

4 Setting Up Weaviate Document Stores

Weaviate serves as the document store for all three vector databases, using separate classes for CoinMarketCap, Binance, and sentiment data. Below is a detailed setup process.

4.1 Weaviate Installation and Configuration

- Install Weaviate locally or on a cloud provider (e.g., AWS).
- Use Docker for local setup:

```
1 docker run -d -p 8080:8080 --name weaviate semitechnologies/weaviate:latest
```

- Configure Weaviate with authentication (optional) and enable hybrid search.
- Set embedding_dim=384 for sentence-transformers/all-MiniLM-L6-v2.

4.2 Vector Database 1: CoinMarketCap API Data

Purpose: Store market data for cryptocurrencies.

Schema:

```
1 {
2   "class": "CoinMarketCap",
3   "properties": [
4     {"name": "crypto_id", "dataType": ["string"]},
5     {"name": "symbol", "dataType": ["string"]},
6     {"name": "price", "dataType": ["number"]},
7     {"name": "market_cap", "dataType": ["number"]},
8     {"name": "volume_24h", "dataType": ["number"]},
9     {"name": "timestamp", "dataType": ["date"]},
10    {"name": "embedding", "dataType": ["number[]"]}
11  ],
12  "vectorizer": "none"
13 }
```

Ingestion Code:

```
1 import requests
2 from haystack_integrations.document_stores.weaviate import
   WeaviateDocumentStore
3 from sentence_transformers import SentenceTransformer
4 from weaviate import Client
5
6 # Initialize Weaviate
7 client = Client("http://localhost:8080")
8 document_store = WeaviateDocumentStore(url="http://localhost:8080",
   embedding_dim=384)
9 model = SentenceTransformer('all-MiniLM-L6-v2')
10
11 # Create schema
12 client.schema.create_class({
13   "class": "CoinMarketCap",
14   "properties": [
15     {"name": "crypto_id", "dataType": ["string"]},
16     {"name": "symbol", "dataType": ["string"]},
17     {"name": "price", "dataType": ["number"]},
18     {"name": "market_cap", "dataType": ["number"]},
19     {"name": "volume_24h", "dataType": ["number"]},
20     {"name": "timestamp", "dataType": ["date"]},
21     {"name": "embedding", "dataType": ["number[]"]}
22   ],
```

```

23     "vectorizer": "none"
24 })
25
26 # Fetch and store data
27 api_key = "YOUR_CMC_API_KEY"
28 url = "https://pro-api.coinmarketcap.com/v1/cryptocurrency/listings/latest"
29 headers = {"X-CMC_PRO_API_KEY": api_key}
30 response = requests.get(url, headers=headers).json()
31
32 documents = []
33 for crypto in response['data']:
34     data = {
35         "crypto_id": str(crypto['id']),
36         "symbol": crypto['symbol'],
37         "price": crypto['quote']['USD']['price'],
38         "market_cap": crypto['quote']['USD']['market_cap'],
39         "volume_24h": crypto['quote']['USD']['volume_24h'],
40         "timestamp": crypto['last_updated']
41     }
42     embedding = model.encode(str(data)).tolist()
43     documents.append({"content": str(data), "meta": data, "embedding":
44                     embedding})
45 document_store.write_documents(documents)

```

4.3 Vector Database 2: Binance API Data with Technical Analysis

Purpose: Store trading data with technical indicators.

Schema:

```

1 {
2     "class": "Binance",
3     "properties": [
4         {"name": "pair", "dataType": ["string"]},
5         {"name": "timestamp", "dataType": ["date"]},
6         {"name": "open", "dataType": ["number"]},
7         {"name": "high", "dataType": ["number"]},
8         {"name": "low", "dataType": ["number"]},
9         {"name": "close", "dataType": ["number"]},
10        {"name": "volume", "dataType": ["number"]},
11        {"name": "rsi", "dataType": ["number"]},
12        {"name": "macd", "dataType": ["number"]},
13        {"name": "bb_upper", "dataType": ["number"]},
14        {"name": "bb_lower", "dataType": ["number"]},
15        {"name": "sto_k", "dataType": ["number"]},
16        {"name": "cci", "dataType": ["number"]},
17        {"name": "atr", "dataType": ["number"]},
18        {"name": "obv", "dataType": ["number"]},
19        {"name": "vwap", "dataType": ["number"]},
20        {"name": "fib_618", "dataType": ["number"]},
21        {"name": "pivot_point", "dataType": ["number"]},
22        {"name": "embedding", "dataType": ["number[]"]}
23    ],
24     "vectorizer": "none"
25 }

```

Ingestion Code:

```

1 import ccxt
2 import pandas as pd
3 import pandas_ta as ta
4 from haystack_integrations.document_stores.weaviate import
    WeaviateDocumentStore
5 from sentence_transformers import SentenceTransformer
6 from weaviate import Client
7
8 # Initialize Weaviate
9 client = Client("http://localhost:8080")
10 document_store = WeaviateDocumentStore(url="http://localhost:8080",
    embedding_dim=384)
11 model = SentenceTransformer('all-MiniLM-L6-v2')
12
13 # Create schema
14 client.schema.create_class({
15     "class": "Binance",
16     "properties": [
17         {"name": "pair", "dataType": ["string"]},
18         {"name": "timestamp", "dataType": ["date"]},
19         {"name": "open", "dataType": ["number"]},
20         {"name": "high", "dataType": ["number"]},
21         {"name": "low", "dataType": ["number"]},
22         {"name": "close", "dataType": ["number"]},
23         {"name": "volume", "dataType": ["number"]},
24         {"name": "rsi", "dataType": ["number"]},
25         {"name": "macd", "dataType": ["number"]},
26         {"name": "bb_upper", "dataType": ["number"]},
27         {"name": "bb_lower", "dataType": ["number"]},
28         {"name": "sto_k", "dataType": ["number"]},
29         {"name": "cci", "dataType": ["number"]},
30         {"name": "atr", "dataType": ["number"]},
31         {"name": "obv", "dataType": ["number"]},
32         {"name": "vwap", "dataType": ["number"]},
33         {"name": "fib_618", "dataType": ["number"]},
34         {"name": "pivot_point", "dataType": ["number"]},
35         {"name": "embedding", "dataType": ["number[]"]}
36     ],
37     "vectorizer": "none"
38 })
39
40 # Fetch Binance data
41 binance = ccxt.binance()
42 ohlcv = binance.fetch_ohlcv('BTC/USDT', timeframe='15m', limit=100)
43 df = pd.DataFrame(ohlcv, columns=['timestamp', 'open', 'high', 'low', '
    close', 'volume'])
44
45 # Compute indicators
46 df['rsi'] = ta.rsi(df['close'], length=14)
47 df['macd'], _, _ = ta.macd(df['close'])
48 df['bb_upper'], _, df['bb_lower'] = ta.bbands(df['close'], length=20)
49 df['sto_k'], _ = ta.stoch(df['high'], df['low'], df['close'])
50 df['cci'] = ta.cci(df['high'], df['low'], df['close'])
51 df['atr'] = ta.atr(df['high'], df['low'], df['close'])
52 df['obv'] = ta.obv(df['close'], df['volume'])
53 df['vwap'] = ta.vwap(df['high'], df['low'], df['close'], df['volume'])
54 df['fib_618'] = df['close'].rolling(20).max() - (df['close'].rolling(20).

```

```

max() - df['close'].rolling(20).min()) * 0.618
55 df['pivot_point'] = (df['high'].shift(1) + df['low'].shift(1) + df['close'
    ].shift(1)) / 3
56
57 # Store in Weaviate
58 documents = []
59 for idx, row in df.iterrows():
60     data = row.to_dict()
61     data['pair'] = 'BTC/USDT'
62     embedding = model.encode(str(data)).tolist()
63     documents.append({"content": str(data), "meta": data, "embedding":
        embedding})
64 document_store.write_documents(documents)

```

4.4 Vector Database 3: News and Social Media Sentiment

Purpose: Store news and X posts with sentiment scores.

Schema:

```

1 {
2   "class": "Sentiment",
3   "properties": [
4     {"name": "source", "dataType": ["string"]},
5     {"name": "text", "dataType": ["text"]},
6     {"name": "crypto_mentioned", "dataType": ["string"]},
7     {"name": "sentiment_score", "dataType": ["number"]},
8     {"name": "timestamp", "dataType": ["date"]},
9     {"name": "embedding", "dataType": ["number[]"]}
10  ],
11  "vectorizer": "none"
12 }

```

Ingestion Code:

```

1 from transformers import pipeline
2 from haystack_integrations.document_stores.weaviate import
    WeaviateDocumentStore
3 from sentence_transformers import SentenceTransformer
4 from weaviate import Client
5 import tweepy
6
7 # Initialize Weaviate
8 client = Client("http://localhost:8080")
9 document_store = WeaviateDocumentStore(url="http://localhost:8080",
    embedding_dim=384)
10 model = SentenceTransformer('all-MiniLM-L6-v2')
11 sentiment_analyzer = pipeline("sentiment-analysis", model="mrms8488/
    distilroberta-finetuned-financial-news-sentiment-analysis")
12
13 # Create schema
14 client.schema.create_class({
15     "class": "Sentiment",
16     "properties": [
17         {"name": "source", "dataType": ["string"]},
18         {"name": "text", "dataType": ["text"]},
19         {"name": "crypto_mentioned", "dataType": ["string"]},
20         {"name": "sentiment_score", "dataType": ["number"]},

```



```

21         {"name": "timestamp", "dataType": ["date"]},
22         {"name": "embedding", "dataType": ["number[]"]}
23     ],
24     "vectorizer": "none"
25 })
26
27 # Fetch X posts
28 client = tweepy.Client(bearer_token="YOUR_X_BEARER_TOKEN")
29 tweets = client.search_recent_tweets(query="#BTC", max_results=100).data
30
31 # Store in Weaviate
32 documents = []
33 for tweet in tweets:
34     text = tweet.text
35     sentiment = sentiment_analyzer(text)[0]
36     data = {
37         "source": "X",
38         "text": text,
39         "crypto_mentioned": "BTC",
40         "sentiment_score": sentiment['score'] if sentiment['label'] == '
positive' else -sentiment['score'],
41         "timestamp": tweet.created_at.isoformat()
42     }
43     embedding = model.encode(text).tolist()
44     documents.append({"content": text, "meta": data, "embedding": embedding
    })
45 document_store.write_documents(documents)

```

4.5 Robustness Recommendations

- Use Weaviate's **hybrid search** for optimal retrieval.
- Implement **data deduplication** using Weaviate's UUIDs.
- Validate API responses with schema checks.
- Schedule **incremental updates** with a cron job.

5 Building the Haystack RAG Pipeline

The Haystack RAG pipeline retrieves data from Weaviate and generates trading strategies using OpenRouter.ai.

5.1 Haystack Components

- **Document Store:** Weaviate with three classes.
- **Embedder:** sentence-transformers/all-MiniLM-L6-v2.
- **Retriever:** WeaviateHybridRetriever with top-k=5.
- **Generator:** OpenRouter.ai's LLaMA 3.1 70B via OpenAIGenerator.
- **Pipeline:** Combines retriever and generator with a custom prompt.

5.2 Pipeline Setup

- **Retriever:** Queries all Weaviate classes for relevant data.
- **Generator:** Synthesizes trading signals based on the confluence strategy.
- **Prompt Template:** Incorporates technical indicators and sentiment.

Pipeline Code:

```
1 from haystack import Pipeline
2 from haystack_integrations.components.retrievers.weaviate import
  WeaviateHybridRetriever
3 from haystack.components.generators import OpenAIGenerator
4 from haystack_integrations.document_stores.weaviate import
  WeaviateDocumentStore
5
6 # Initialize components
7 document_store = WeaviateDocumentStore(url="http://localhost:8080",
  embedding_dim=384)
8 retriever = WeaviateHybridRetriever(document_store=document_store, top_k=5)
9 generator = OpenAIGenerator(
10     api_key="YOUR_OPENROUTER_API_KEY",
11     api_base_url="https://openrouter.ai/api/v1",
12     model="meta-llama/llama-3.1-70b-instruct"
13 )
14
15 # Create pipeline
16 pipeline = Pipeline()
17 pipeline.add_component("retriever", retriever)
18 pipeline.add_component("generator", generator)
19 pipeline.connect("retriever.documents", "generator.documents")
20
21 # Query example
22 query = "Find trading opportunities for BTC/USDT with oversold RSI and
  positive sentiment"
23 result = pipeline.run({
24     "retriever": {"query": query},
25     "generator": {
26         "prompt": """
27         Given the following data:
28         {documents}
29         Suggest a trading strategy for BTC/USDT using:
30         - 50 EMA > 200 EMA for trend
31         - RSI crossing above 30
32         - Price near 0.618 Fibonacci support
33         - MACD bullish crossover
34         - Volume spike (OBV increasing)
35         Set stop-loss at 2x ATR below entry and take-profit at 2x risk.
36         """
37     }
38 })
39 print(result["generator"]["replies"])
```

6 Developing the Haystack Agent

The Haystack Agent implements the confluence strategy and executes trades.

6.1 Agent Setup

- **Tools:** Retrieve data, compute indicators, execute trades.
- **Logic:** Combines RAG outputs with strategy rules.
- **Execution:** Uses Binance API with risk management.

Agent Code:

```
1 from haystack import Pipeline
2 from haystack.components.generators import OpenAIGenerator
3 from haystack.agents import Agent, Tool
4 import ccxt
5 import pandas as pd
6 import pandas_ta as ta
7
8 # Trade execution tool
9 def place_trade(symbol, action, price, amount, stop_loss, take_profit):
10     binance = ccxt.binance({'apiKey': 'YOUR_API_KEY', 'secret': 'YOUR_API_SECRET'})
11     try:
12         if action == "buy":
13             order = binance.create_limit_buy_order(symbol, amount, price)
14             # Set stop-loss and take-profit (simplified)
15             return f"Placed buy order: {order}"
16         elif action == "sell":
17             order = binance.create_limit_sell_order(symbol, amount, price)
18             return f"Placed sell order: {order}"
19     except Exception as e:
20         return f"Error: {e}"
21
22 trade_tool = Tool(
23     name="TradeExecutor",
24     description="Executes trades on Binance with stop-loss and take-profit",
25     function=place_trade
26 )
27
28 # Indicator computation tool
29 def compute_indicators(symbol, timeframe):
30     binance = ccxt.binance()
31     ohlcv = binance.fetch_ohlcv(symbol, timeframe, limit=100)
32     df = pd.DataFrame(ohlcv, columns=['timestamp', 'open', 'high', 'low', 'close', 'volume'])
33     df['ema50'] = ta.ema(df['close'], length=50)
34     df['ema200'] = ta.ema(df['close'], length=200)
35     df['rsi'] = ta.rsi(df['close'], length=14)
36     df['macd'], _, _ = ta.macd(df['close'])
37     df['obv'] = ta.obv(df['close'], df['volume'])
38     df['atr'] = ta.atr(df['high'], df['low'], df['close'])
39     return df.iloc[-1].to_dict()
40
41 indicator_tool = Tool(
42     name="IndicatorCalculator",
43     description="Computes technical indicators for a trading pair",
44     function=compute_indicators
45 )
46
47 # Initialize agent
```

```

48 generator = OpenAIGenerator(
49     api_key="YOUR_OPENROUTER_API_KEY",
50     api_base_url="https://openrouter.ai/api/v1",
51     model="meta-llama/llama-3.1-70b-instruct"
52 )
53 agent = Agent(generator=generator, tools=[trade_tool, indicator_tool])
54
55 # Run agent
56 result = agent.run(
57     query="For BTC/USDT on 15m timeframe, check if 50 EMA > 200 EMA, RSI > 30, and positive sentiment. If true, buy 0.01 BTC at market price with 2x ATR stop-loss.",
58     prompt="Use IndicatorCalculator to check conditions and TradeExecutor to place the trade if valid."
59 )
60 print(result)

```

7 Building the Frontend Application

Tech Stack:

- Framework: React with Tailwind CSS.
- Visualization: Chart.js for candlestick charts and indicators.
- Backend: FastAPI for Haystack integration.

Features:

- Configure trading pairs, risk levels, and strategy parameters.
- Display real-time price charts, indicators, and sentiment.
- Approve or automate trades.
- Backtest strategies on historical data.

Frontend Code (React):

```

1 import React, { useState, useEffect } from 'react';
2 import { Line } from 'react-chartjs-2';
3 import axios from 'axios';
4
5 const Dashboard = () => {
6     const [priceData, setPriceData] = useState([]);
7     const [indicators, setIndicators] = useState({});
8
9     useEffect(() => {
10         // Fetch price and indicator data
11         axios.get('/api/price/BTCUSDT').then(response => {
12             setPriceData(response.data);
13         });
14         axios.get('/api/indicators/BTCUSDT').then(response => {
15             setIndicators(response.data);
16         });
17     }, []);
18
19     const chartData = {
20         labels: priceData.map(d => d.timestamp),

```

```

21     datasets: [
22       { label: 'Price', data: priceData.map(d => d.close), borderColor: '
          blue' },
23       { label: 'EMA 50', data: priceData.map(d => d.ema50), borderColor: '
          green' },
24       { label: 'EMA 200', data: priceData.map(d => d.ema200), borderColor:
          'red' }
25     ]
26   };
27
28   return (
29     <div className="p-4">
30       <h2 className="text-2xl font-bold">BTC/USDT Dashboard</h2>
31       <Line data={chartData} />
32       <div>RSI: {indicators.rsi?.toFixed(2)}</div>
33       <div>MACD: {indicators.macd?.toFixed(2)}</div>
34     </div>
35   );
36 };
37
38 export default Dashboard;

```

8 Integration and Testing

Integration:

- Connect Haystack pipeline and agent to FastAPI.
- Use Binance testnet API for trade testing.

Testing:

- Validate data ingestion with schema checks.
- Test retriever accuracy with queries like “RSI < 30 and sentiment > 0.7”.
- Simulate trades to ensure strategy compliance.

9 Deployment

9.1 Deployment Setup

- **Weaviate:** Deploy on AWS EC2 using Docker.
- **Haystack Backend:** Deploy FastAPI on AWS ECS with Docker.
- **Frontend:** Host React app on AWS Amplify with Cloudflare CDN.
- **Message Queue:** Use Kafka on AWS MSK for data ingestion.

9.2 Deployment Steps

1. Provision Infrastructure:

- Launch EC2 instance for Weaviate (e.g., t3.medium).
- Set up ECS cluster for FastAPI (use Fargate for serverless).

- Configure Amplify for React frontend.
 - Deploy Kafka on MSK for data pipelines.
2. **Configure Weaviate:**
 - Run Weaviate Docker container with persistent storage (EBS).
 - Set environment variables for authentication and hybrid search.
 3. **Deploy Haystack Backend:**
 - Create Docker image for FastAPI with Haystack dependencies.
 - Push to AWS ECR and deploy to ECS.
 - Configure load balancer (ALB) for API endpoints.
 4. **Deploy Frontend:**
 - Build React app and push to Amplify.
 - Configure Cloudflare for CDN and DDoS protection.
 5. **Set Up Monitoring:**
 - Use Prometheus and Grafana on EC2 for metrics.
 - Monitor API latency, trade execution, and Weaviate performance.
 6. **Security:**
 - Encrypt API keys with AWS KMS.
 - Use IAM roles for ECS and Amplify.
 - Enable WAF on Cloudflare.

9.3 Maintenance

- Schedule Weaviate backups using EBS snapshots.
- Monitor OpenRouter.ai and Binance API rate limits.
- Update Haystack and LLM models regularly.

10 Advanced Features

- **On-Chain Integration:** Use Etherscan for wallet movements.
- **Backtesting:** Simulate strategies with historical Binance data.
- **Explainability:** Display reasons for agent decisions.

11 Regulatory Compliance

- Ensure KYC/AML compliance for trade execution.
- Include risk disclaimers in the frontend.
- Redirect to <https://help.x.com/en/using-x/x-premium> for subscriptions and <https://x.ai/api> for API pricing.

12 Conclusion

This guide provides a robust framework for building an Agentic RAG model for crypto trading using Haystack, Weaviate, and OpenRouter.ai. The confluence-based trading strategy ensures high-confidence decisions, while detailed deployment steps enable production readiness.

Next Steps:

1. Deploy Weaviate and set up data ingestion.
2. Build and test the Haystack pipeline.
3. Develop the React frontend.
4. Deploy on AWS and monitor performance.